

Dependency modelling toolbox

Leo Lahti*, Olli-Pekka Huovilainen, Abhishek Tripathi,
Ilkka Huopaniemi, and Tommi Suvitaival
Department of Information and Computer Science,
Aalto University School of Science and Technology, Finland

June 23, 2010

1 Introduction

Investigation of dependencies between multiple data sources allows the discovery of regularities and interactions that are not seen in individual data sets. The importance of such methods is increasing with the availability and size of co-occurring data sets in computational biology, open data initiatives, and in other domains. Practical, open access implementations of general-purpose algorithms will help to realize the full potential of these information sources.

This package provides general-purpose tools for the discovery and analysis of statistical dependencies between co-occurring data sources. The implementations are based on well-established models such as probabilistic canonical correlation analysis and multi-task learning [1, 2, 3, 4, 5]. Probabilistic framework deals rigorously with the uncertainties associated with small sample sizes, and allows incorporation of prior information in the analysis through Bayesian priors [4]. The applicability of the models has been demonstrated in previous case studies [3, 4, 5]. This is a development version. Your feedback and contributions are welcome. See the project page at R-Forge¹, or contact project authors.

2 Available functionality

Current version is divided in three independent modules. We are working on integrating these application-oriented modules into a unified dependency modelling toolbox. The functionality includes

- regularized dependency detection (pint²; [4])
- dependency-based dimensionality reduction (drCCA³; [5])
- multi-way modeling of co-occurrence data (multiWayCCA⁴; [3])

*leo.lahti@iki.fi

¹<http://dmt.r-forge.r-project.org/>

²<http://bioconductor.org/packages/release/bioc/html/pint.html>

³<http://www.cis.hut.fi/projects/mi/software/drCCA/>

⁴<http://www.cis.hut.fi/projects/mi/software/multiWayCCA/>

Below is a brief summary of each module together with installation instructions.

3 Probabilistic dependency modeling framework (pint)




Figure 1: Graphical description of the shared latent variable model showing generation of data sets X and Y from latent shared variable \mathbf{z} through W_x and W_y

The pint package⁵; [4] implements the probabilistic dependency modeling framework presented in [2] and its extensions [1, 6, 4]. This is a latent variable model that assumes that the two data sets, X and Y can be decomposed in *shared* and *data set-specific* components (Figure~1). The model helps to discover the shared components, given modeling assumptions.

The shared signal of two data sets is modeled with a shared latent variable \mathbf{z} . Intuitively, this measures the strength of the shared signal in each patient. While the variation is shared, it can have different manifestation in each data set. This is described by $W_x \mathbf{z}$ and $W_y \mathbf{z}$ where W_x, W_y indicate how the shared signal is observed in the individual data sets. Assuming a standard Gaussian model for the shared latent variable $\mathbf{z} \sim \mathcal{N}(0, I)$ and data set-specific effects, this leads to the following model:

$$\begin{aligned}
 X &\sim W_x \mathbf{z} + \varepsilon_x \\
 Y &\sim W_y \mathbf{z} + \varepsilon_y \\
 \varepsilon_{\cdot} &\sim \mathcal{N}(0, \Psi_{\cdot}) \\
 \mathbf{z} &\sim \mathcal{N}(0, I)
 \end{aligned} \tag{1}$$

The data set-specific effects are described by the covariance matrices Ψ_x, Ψ_y . The model parameters are estimated with an expectation-maximization (EM) algorithm.

3.1 Special cases

Special cases of the model include probabilistic versions of canonical correlation analysis, factor analysis, and principal component analysis, and regularized versions of them.

⁵<http://bioconductor.org/packages/release/bioc/html/pint.html>

Probabilistic CCA (pCCA) assumes full covariance matrices Ψ_x, Ψ_y . This gives the most detailed model for the data set specific effects. The connection of this latent variable model and the traditional canonical correlation analysis has been established in [2].

Probabilistic factor analysis (pFA) is obtained with a diagonal covariances Ψ_x, Ψ_y . In addition, a special case is implemented where each covariance matrix Ψ is isotropic but they are not necessarily identical (as would be the case in pPCA). This model is identical to concatenating X, Y , and fitting ordinary probabilistic factor analysis on the concatenated data set. The structure of the covariances is simpler than in pCCA. This regularizes the solution and can potentially reduce overfitting in some applications.

Probabilistic PCA (pPCA) is obtained with identical isotropic covariances for the data set-specific effects: $\Psi_x = \Psi_y = \sigma I$. This model is identical to concatenating X, Y , and fitting ordinary probabilistic PCA on the concatenated data set.

3.2 Regularized dependency modeling

We provide tools to guide dependency modeling through Bayesian priors [4]. Similarity-constrained probabilistic CCA (pSimCCA) imposes a prior on the relation between W_x and W_y . This can be used to guide modeling to focus on certain types of dependencies, and to avoid overfitting. The relationship is described through $W_y = TW_x$. A prior on T can be used to focus the modeling on certain types of dependencies. We use matrix normal prior distribution: $P(T) = N_m(H, \sigma_T^2 I, \sigma_T^2 I)$. By default, $H = I$ and $\sigma_T^2 = 0$, which results in identical manifestation of the shared signal in the two data sets: $W_y = W_x$. This model is denoted pSimCCA in the package. However, the prior can be loosened by tuning σ_T^2 . With $\sigma_T^2 \rightarrow \infty$, estimation of W_x and W_y become independent, which leads to ordinary probabilistic CCA. It is also possible to tune the mean matrix H . This would set a particular relationship between the manifestations of the shared component in each data set, and σ_T^2 is again be used to tune the strength of such prior.

3.3 Installing pint

Install pint from within R with command
`source('http://bioconductor.org/biocLite.R')`
`biocLite('pint')`

3.4 Documentation of pint

The package implements the dependency modeling framework explained above (see function `'fit.dependency.model'`), and provides wrappers for the special cases of the model. The documentation is available at <http://bioconductor.org/packages/2.6/bioc/vignettes/pint/inst/doc/depsearch.pdf>

3.5 Applications of pint

pint has been used to integrate gene expression and copy number (aCGH) data to discover cancer-associated chromosomal regions. See [4] for details.

4 Dependency-based dimensionality reduction (drCCA)

The drCCA⁶ [5] method retains the variation that is shared between the original data sources, while reducing the dimensionality by ignoring variation that is specific to any of the data sources alone. This captures the common variation that is shared by the measurement sources. Note that the shared variation can be reflected in different ways in the different data sets. The approach utilizes generalized canonical correlation analysis to perform a linear projection on the collection data sets. Linearity makes it fast on large data sets. The package includes regularization and tools for selecting the final dimensionality of the combined data set automatically.

4.1 Installing drCCA

Install drCCA from within R with command
`'install.packages("drCCA", repos="http://R-Forge.R-project.org")'`

4.2 Documentation of drCCA

Currently only online-documentation for the package is available. See <http://www.cis.hut.fi/projects/mi/software/drCCA/dochtml/00Index.html>

4.3 Applications of drCCA

drCCA has been applied to dimensionality reduction in functional genomics [5]. It combines information from several measurement sources. This helps to reduce the noise in biological experiments with high dimensionality but small sample size. The drCCA can be used to summarize shared variation in several data sources with co-occurring samples into a one vectorial data set of lower dimensionality which captures the shared effects of the data sources.

5 Multi-way multi-view models (multiWayCCA)

multiWayCCA⁷ provides tools for multi-way, multi-source modeling. This is particularly useful for simultaneous multi-way (anova-type) modelling of multiple related data sources. For details, see the original paper [3].

5.1 Installing multiWayCCA

multiWayCCA is currently available only as example source code. The sources contain both code and examples. Download the source⁸. Then uncompress the folder; `readme.txt` in the uncompressed folder contains instructions for running the analysis. For application examples, see [3].

⁶<http://www.cis.hut.fi/projects/mi/software/drCCA/>

⁷<http://www.cis.hut.fi/projects/mi/software/multiWayCCA/>

⁸<http://www.cis.hut.fi/projects/mi/software/multiWayCCA/multiWayCCA-package-100326.zip>

5.2 Documentation of multiWayCCA

For documentation and examples, see the readme.txt file included in the package.

5.3 Applications of multiWayCCA: metabolomics

multiWayCCA has been applied in high-throughput metabolomics studies, see [3] for details.

5.4 Licensing terms

Currently the different modules have different licensing terms. The licenses will be unified in the integrated version.

- pint: GNU GPL ≥ 2
- drCCA: GNU LGPL ≥ 2.1
- multiWayCCA: no license associated with the code at the moment

Acknowledgements

The project is a joint effort by several people: Leo Lahti, Tommi Suvitaival, Olli-Pekka Huovilainen, Ilkka Huopaniemi, Abhishek Tripathi, Arto Klami, and Samuel Kaski from the Statistical Machine Learning and Bioinformatics group at the Department of Information and Computer Science, Aalto University School of Science and Technology, Finland. Abhishek Tripathi is with Department of Computer Science, University of Helsinki, Finland. The authors belong also to Helsinki Institute for Information Technology HIIT and Adaptive Informatics Research Centre AIRC.

References

- [1] Cédric Archambeau, Nicolas Delannay, and Michel Verleysen. Robust probabilistic projections. In W.W. Cohen and A. Moore, editors, *Proceedings of the 23rd International conference on machine learning*, pages 33–40. ACM, 2006.
- [2] Francis R. Bach and Michael I. Jordan. A probabilistic interpretation of canonical correlation analysis. Technical Report 688, Department of Statistics, University of California, Berkeley, 2005.
- [3] Ilkka Huopaniemi, Tommi Suvitaival, Janne Nikkilä, Matej Orešić, and Samuel Kaski. Multivariate multi-way analysis of multi-source data. *Bioinformatics*, 2010. (ISMB 2010, to appear).
- [4] Leo Lahti, Samuel Myllykangas, Sakari Knuutila, and Samuel Kaski. Dependency detection with similarity constraints. In *Proc. MLSP'09 IEEE International Workshop on Machine Learning for Signal Processing*, 2009.

- [5] Arto Klami, Abhishek Tripathi and Samuel Kaski. Simple integrative pre-processing preserves what is shared in data sources. *BMC Bioinformatics*, 9(111), 2008.
- [6] Arto Klami and Samuel Kaski. Probabilistic approach to detecting dependencies between data sets. *Neurocomputing*, 72(1-3):39–46, 2008.